

# Numerical Optimization

---

Joachim Vandekerckhove

# Motivating Problems for Numerical Optimization

- **Function Minimization:** Find the minimum of a function  $f(x)$  over a given range of  $x$ .

# Motivating Problems for Numerical Optimization

- **Function Minimization:** Find the minimum of a function  $f(x)$  over a given range of  $x$ .
- **Parameter Estimation:** Given a set of observations, estimate the parameters of a model that best fit the data.

# Motivating Problems for Numerical Optimization

- **Function Minimization:** Find the minimum of a function  $f(x)$  over a given range of  $x$ .
- **Parameter Estimation:** Given a set of observations, estimate the parameters of a model that best fit the data.
- **Machine Learning:** Train a model to predict outcomes based on input features by minimizing a loss function.

# Motivating Problems for Numerical Optimization

- **Function Minimization:** Find the minimum of a function  $f(x)$  over a given range of  $x$ .
- **Parameter Estimation:** Given a set of observations, estimate the parameters of a model that best fit the data.
- **Machine Learning:** Train a model to predict outcomes based on input features by minimizing a loss function.
- **Procedure Optimization:** Find the conditions where a certain procedure works best.

# Numerical optimization in cognitive modeling

- Cognitive models have parameters that need to be estimated

# Numerical optimization in cognitive modeling

- Cognitive models have parameters that need to be estimated
- Estimation means minimizing the **loss**  $\ell(x, \theta)$ ...

$$\hat{\theta} = \arg \min_{\theta} (\ell(x; \theta))$$

# Numerical optimization in cognitive modeling

- Cognitive models have parameters that need to be estimated
- Estimation means minimizing the **loss**  $\ell(x, \theta)$ ...

$$\hat{\theta} = \arg \min_{\theta} (\ell(x; \theta))$$

- ... or maximizing the **likelihood**  $f(x|\theta)$ :

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} (f(x|\theta)) \\ &= \arg \min_{\theta} (-\log (f(x|\theta)))\end{aligned}$$



# Numerical optimization in cognitive modeling

- Cognitive models have parameters that need to be estimated
- Estimation means minimizing the **loss**  $\ell(\mathbf{x}, \theta)$ ...

$$\hat{\theta} = \arg \min_{\theta} (\ell(\mathbf{x}; \theta))$$

- ... or maximizing the **likelihood**  $f(\mathbf{x}|\theta)$ :

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} (f(\mathbf{x}|\theta)) \\ &= \arg \min_{\theta} (-\log (f(\mathbf{x}|\theta)))\end{aligned}$$

- To-be-optimized function is called target or **objective** function

# Numerical optimization in cognitive modeling

Optimization is sometimes possible analytically but most often done using numerical methods:

1. Guess  $\theta$  through some method

# Numerical optimization in cognitive modeling

Optimization is sometimes possible analytically but most often done using numerical methods:

1. Guess  $\theta$  through some method
2. Compute objective function

# Numerical optimization in cognitive modeling

Optimization is sometimes possible analytically but most often done using numerical methods:

1. Guess  $\theta$  through some method
2. Compute objective function
3. Check if the current guess is best yet

# Numerical optimization in cognitive modeling

Optimization is sometimes possible analytically but most often done using numerical methods:

1. Guess  $\theta$  through some method
2. Compute objective function
3. Check if the current guess is best yet
4. Repeat until good enough

# Numerical optimization in cognitive modeling

---

**Algorithm:** Abstract Numerical Minimization

---

**Input** : objective function  $\ell(\theta)$

**Output:** the learned parameters  $\theta$

1. Initial guess parameters  $\theta$
  2. **while**  $\ell(\theta)$  is too high **do**
    - └ Propose new parameters:  $\theta \leftarrow \theta + \text{update}$ ;
  3. **return**  $\theta$ ;
-

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum



# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:
    1. Guess  $\theta_0$  through some method and compute  $\ell(x; \theta_0)$

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:
    1. Guess  $\theta_0$  through some method and compute  $\ell(x; \theta_0)$
    2. Guess  $\theta_1$  nearby and compute  $\ell(x; \theta_1)$

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:
    1. Guess  $\theta_0$  through some method and compute  $\ell(x; \theta_0)$
    2. Guess  $\theta_1$  nearby and compute  $\ell(x; \theta_1)$
    3. If  $\ell(x; \theta_1) < \ell(x; \theta_0)$ , move further past  $\theta_1$ , otherwise move back beyond  $\theta_0$

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:
    1. Guess  $\theta_0$  through some method and compute  $\ell(x; \theta_0)$
    2. Guess  $\theta_1$  nearby and compute  $\ell(x; \theta_1)$
    3. If  $\ell(x; \theta_1) < \ell(x; \theta_0)$ , move further past  $\theta_1$ , otherwise move back beyond  $\theta_0$
    4. Decrease step size a little

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:
    1. Guess  $\theta_0$  through some method and compute  $\ell(x; \theta_0)$
    2. Guess  $\theta_1$  nearby and compute  $\ell(x; \theta_1)$
    3. If  $\ell(x; \theta_1) < \ell(x; \theta_0)$ , move further past  $\theta_1$ , otherwise move back beyond  $\theta_0$
    4. Decrease step size a little
    5. Repeat until good enough

# Numerical optimization in cognitive modeling

- Choice of numerical method depends on the problem and available resources, there's **no free lunch**
- **Gradient-based optimization** methods use derivatives to search for the minimum
  - Derivatives can be approximated using finite differences:
    1. Guess  $\theta_0$  through some method and compute  $\ell(x; \theta_0)$
    2. Guess  $\theta_1$  nearby and compute  $\ell(x; \theta_1)$
    3. If  $\ell(x; \theta_1) < \ell(x; \theta_0)$ , move further past  $\theta_1$ , otherwise move back beyond  $\theta_0$
    4. Decrease step size a little
    5. Repeat until good enough
- Other optimization methods include genetic algorithms, simulated annealing, particle swarm optimization, and the **Nelder-Mead simplex**

# Gradient-based optimization



# Gradient-based optimization

## Gradient-based optimization

- Gradient descent works by iteratively adjusting the model's parameters in the direction of steepest descent of the objective function.

# Gradient-based optimization


- Gradient descent works by iteratively adjusting the model's parameters in the direction of steepest descent of the objective function.
- Newton's method uses the **Hessian matrix** (i.e., the local curvature) to adjust the parameters.

# Gradient-based optimization

- Gradient descent works by iteratively adjusting the model's parameters in the direction of steepest descent of the objective function.
- Newton's method uses the **Hessian matrix** (i.e., the local curvature) to adjust the parameters.
- Quasi-Newton methods approximate the Hessian matrix using a limited amount of information.


# Gradient-based optimization

- Gradient descent works by iteratively adjusting the model's parameters in the direction of steepest descent of the objective function.
- Newton's method uses the **Hessian matrix** (i.e., the local curvature) to adjust the parameters.
- Quasi-Newton methods approximate the Hessian matrix using a limited amount of information.



# Gradient-based optimization

- Gradient descent works by iteratively adjusting the model's parameters in the direction of steepest descent of the objective function.
- Newton's method uses the **Hessian matrix** (i.e., the local curvature) to adjust the parameters.
- Quasi-Newton methods approximate the Hessian matrix using a limited amount of information.
- Some nice visualizations:  
<https://www.benfrederickson.com/numerical-optimization/>



# Simulated Annealing

- Simulated annealing is a powerful **global optimization method** that is inspired by the process of annealing in metallurgy.

# Simulated Annealing

- Simulated annealing is a powerful **global optimization method** that is inspired by the process of annealing in metallurgy.
- The process is governed by a **temperature** setting that determines the probability of accepting a step in the wrong direction.



# Simulated Annealing

- Simulated annealing is a powerful **global optimization method** that is inspired by the process of annealing in metallurgy.
- The process is governed by a **temperature** setting that determines the probability of accepting a step in the wrong direction.
- A hot algorithm will move randomly, with little regard for the objective function value. A cold algorithm performs steepest ascent.

# Simulated Annealing

- Simulated annealing is a powerful **global optimization method** that is inspired by the process of annealing in metallurgy.
- The process is governed by a **temperature** setting that determines the probability of accepting a step in the wrong direction.
- A hot algorithm will move randomly, with little regard for the objective function value. A cold algorithm performs steepest ascent.
- A slow cooling schedule allows the algorithm to explore many possible peaks before climbing to the top.

# Simulated Annealing

---

**Algorithm 2:** Simulated Annealing Algorithm

---

**Input:** Initial guess  $s$ , initial temperature  $T$ ; objective  $\ell(\theta)$

**Output:** Best parameter set found during search

**while** stopping condition not met **do**

    Generate new guess  $s'$  by random perturbation to  $s$ ;

    Calculate loss difference  $\Delta\ell = \ell(s') - \ell(s)$ ;

**if**  $\Delta E < 0$  **then**

        Accept  $s'$  as the new current state  $s$ ;

**else**

        Accept  $s'$  with probability  $\exp(-\Delta\ell/T)$ ;

**end**





    Decrease temperature  $T$ ;

**end**

**return** Best state found during search

---

# Simulated Annealing



# Constrained vs. Unconstrained Optimization

- Two main types of optimization problems: constrained and unconstrained

# Constrained vs. Unconstrained Optimization

- Two main types of optimization problems: constrained and unconstrained
- In an **unconstrained optimization** problem, the goal is to minimize (or maximize) a function subject to no constraints

# Constrained vs. Unconstrained Optimization

- Two main types of optimization problems: constrained and unconstrained
- In an **unconstrained optimization** problem, the goal is to minimize (or maximize) a function subject to no constraints
- In a **constrained optimization** problem, the goal is to minimize (or maximize) a function subject to one or more constraints

# Constrained vs. Unconstrained Optimization

- Two main types of optimization problems: constrained and unconstrained
- In an **unconstrained optimization** problem, the goal is to minimize (or maximize) a function subject to no constraints
- In a **constrained optimization** problem, the goal is to minimize (or maximize) a function subject to one or more constraints
- Constrained optimization often requires the use of specialized optimization algorithms and techniques to find the global minimum (or maximum)



# Constrained vs. Unconstrained Optimization

Constraints can be that some linear transformation  $A$  of the parameters must be less than some constant  $b$ , and/or that some linear transformation  $A_{eq}$  of the parameters must be equal to some constant  $b_{eq}$ , and/or that they must fall in bounds  $(l_b, u_b)$ , and/or that some are integers.

$$\begin{array}{ll} \underset{x}{\text{minimize}} & f^T x \\ \text{subject to} & \left\{ \begin{array}{l} Ax \leq b \text{ (MATLAB)} \\ Ax \geq b \text{ (python)} \\ A_{eq}x = b_{eq} \\ l_b \leq x \leq u_b \\ x(\text{intflag}) \text{ are integers} \end{array} \right\} \end{array}$$

## Unconstrained Example: Production Environment

- Imagine a company wants to optimize its production process by finding the optimal temperature and pressure settings for the production equipment.

## Unconstrained Example: Production Environment

- Imagine a company wants to optimize its production process by finding the optimal temperature and pressure settings for the production equipment.
- The cost of production is given by:

$$\min_x f(x) = e^{(x_1-2)^2+(x_2-0.5)^2}$$

## Unconstrained Example: Production Environment

- Imagine a company wants to optimize its production process by finding the optimal temperature and pressure settings for the production equipment.
- The cost of production is given by:

$$\min_x f(x) = e^{(x_1-2)^2+(x_2-0.5)^2}$$

- What are the optimal settings to run the production process?

## Unconstrained Example: Parameter Estimation

- Assume we have some data  $D \leftrightarrow \{5, -7, 3, 0, -11\}$

## Unconstrained Example: Parameter Estimation

- Assume we have some data  $D \leftrightarrow \{5, -7, 3, 0, -11\}$
- We want to describe  $D$  with a single number such that the squared distance to each data point is minimal:

$$\min_x f(x) = \sum_{i=1}^5 (x - d_i)^2$$

## Unconstrained Example: Parameter Estimation

- Assume we have some data  $D \leftrightarrow \{5, -7, 3, 0, -11\}$
- We want to describe  $D$  with a single number such that the squared distance to each data point is minimal:

$$\min_x f(x) = \sum_{i=1}^5 (x - d_i)^2$$

- What is the best  $x$ ?

## Unconstrained Example: Signal Detection Theory

- Assume we have some data  $D \leftrightarrow \{\text{Hits} = 4, \text{Misses} = 1, \text{False alarms} = 2, \text{Correct rejections} = 3\}$



## Unconstrained Example: Signal Detection Theory

- Assume we have some data  $D \leftrightarrow \{\text{Hits} = 4, \text{Misses} = 1, \text{False alarms} = 2, \text{Correct rejections} = 3\}$
- Which  $(d', c) \leftrightarrow x$  make  $D$  most likely:

$$\begin{aligned} (\max_x) \quad f(x) &= \binom{5}{4} \theta_H^5 (1 - \theta_H)^1 \times \binom{5}{2} \theta_F^2 (1 - \theta_F)^3 \\ &\text{with } \begin{cases} \theta_H &= 1 - \Phi\left(-\frac{1}{2}d' + c\right) \\ \theta_F &= 1 - \Phi\left(\frac{1}{2}d' + c\right) \end{cases} \end{aligned}$$

## Unconstrained Example: Signal Detection Theory

- Assume we have some data  $D \leftrightarrow \{\text{Hits} = 4, \text{Misses} = 1, \text{False alarms} = 2, \text{Correct rejections} = 3\}$
- Which  $(d', c) \leftrightarrow x$  make  $D$  most likely:

$$\begin{aligned} (\max_x) \quad f(x) &= \binom{5}{4} \theta_H^5 (1 - \theta_H)^1 \times \binom{5}{2} \theta_F^2 (1 - \theta_F)^3 \\ &\text{with } \begin{cases} \theta_H &= 1 - \Phi\left(-\frac{1}{2}d' + c\right) \\ \theta_F &= 1 - \Phi\left(\frac{1}{2}d' + c\right) \end{cases} \end{aligned}$$

## Unconstrained Example: Signal Detection Theory

- Assume we have some data  $D \leftrightarrow \{\text{Hits} = 4, \text{Misses} = 1, \text{False alarms} = 2, \text{Correct rejections} = 3\}$
- Which  $(d', c) \leftrightarrow x$  make  $D$  most likely:

$$(\max_x) \quad f(x) = \binom{5}{4} \theta_H^5 (1 - \theta_H)^1 \times \binom{5}{2} \theta_F^2 (1 - \theta_F)^3$$

$$\text{with } \begin{cases} \theta_H = 1 - \Phi\left(-\frac{1}{2}d' + c\right) \\ \theta_F = 1 - \Phi\left(\frac{1}{2}d' + c\right) \end{cases}$$

$$\begin{aligned} \Rightarrow (\min_x) \quad f(x) &= -\log \left[ \theta_H^5 (1 - \theta_H)^1 \times \theta_F^2 (1 - \theta_F)^3 \right] \\ &= -5 \log(\theta_H) - \log(1 - \theta_H) \\ &\quad - 2 \log(\theta_F) - 3 \log(1 - \theta_F) \end{aligned}$$

# MATLAB

---

# Unconstrained Optimization using `fminsearch`

$$\min_x f(x) = e^{(x_1-2)^2 + (x_2-0.5)^2}$$

- Set up the problem as an anonymous function:

```
% Define the objective function  
fcn = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2 );
```

# Unconstrained Optimization using `fminsearch`

$$\min_x f(x) = e^{(x_1-2)^2 + (x_2-0.5)^2}$$

- Set up the problem as an anonymous function:

```
% Define the objective function  
fcn = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2 );
```

- Minimize using `fminsearch`:

```
% Solve the optimization problem  
start = [0, 0];  
x = fminsearch(fcn, start);
```

# Unconstrained Optimization using `fminsearch`

$$\min_x f(x) = e^{(x_1-2)^2 + (x_2-0.5)^2}$$

- Set up the problem as an anonymous function:

```
% Define the objective function  
fcn = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2 );
```

- Minimize using `fminsearch`:

```
% Solve the optimization problem  
start = [0, 0];  
x = fminsearch(fcn, start);
```

- The solution is  $x_1 = 2$  and  $x_2 = 0.5$ .

# Unconstrained Optimization using `fminunc`

```
fcu = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2);  
x0 = [0.5, 1.5];  
options = optimoptions('fminunc', 'Display', 'iter', ...  
                        'Algorithm', 'quasi-newton', ...  
                        'MaxIterations', 1000);  
[x, fval, exitflag, output] = fminunc(fcu, x0, options);
```

- `Display=iter` shows optimization progress at each iteration



# Unconstrained Optimization using `fminunc`

```
fcu = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2);  
x0 = [0.5, 1.5];  
options = optimoptions('fminunc', 'Display', 'iter', ...  
                        'Algorithm', 'quasi-newton', ...  
                        'MaxIterations', 1000);  
[x, fval, exitflag, output] = fminunc(fcu, x0, options);
```

- `Display=iter` shows optimization progress at each iteration
- `Algorithm` specifies the optimization algorithm to use

# Unconstrained Optimization using `fminunc`

```
fcn = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2);  
x0 = [0.5, 1.5];  
options = optimoptions('fminunc', 'Display', 'iter', ...  
                        'Algorithm', 'quasi-newton', ...  
                        'MaxIterations', 1000);  
[x, fval, exitflag, output] = fminunc(fcn, x0, options);
```

- `Display=iter` shows optimization progress at each iteration
- `Algorithm` specifies the optimization algorithm to use
- `MaxIterations` sets the maximum number of iterations

# Unconstrained Optimization using `fminunc`

```
fcu = @(x) exp( (x(1)-2).^2 + (x(2)-.5).^2);  
x0 = [0.5, 1.5];  
options = optimoptions('fminunc', 'Display', 'iter', ...  
                        'Algorithm', 'quasi-newton', ...  
                        'MaxIterations', 1000);  
[x, fval, exitflag, output] = fminunc(fcu, x0, options);
```

- `Display=iter` shows optimization progress at each iteration
- `Algorithm` specifies the optimization algorithm to use
- `MaxIterations` sets the maximum number of iterations
- Output variables `x`, `fval`, `exitflag`, and `output` contain the optimization results and information

## fminunc vs. fmincon

Using fminunc:

```
fcu = @(x)sum((x-[5, -7, 3, 0, -11]).^2);  
x0 = 0.5;  
[x, fval, exitflag, output] = fminunc(fcu, x0);
```

Enforce  $x \geq 0$  with fmincon:

```
fcu = @(x)sum((x-[5, -7, 3, 0, -11]).^2);  
x0 = 0.5;  
[x, fval, exitflag, output] = fmincon(fcu, x0, -1, 0);
```

Note that  $A = -1$  and  $b = 0$ , so we enforce that  $-x \leq 0$

MATLAB has a lot of numerical optimization features.

As of September 2022, the MATLAB Optimization Toolbox User's Guide is 1,584 pages long. It is an excellent review of modern numerical optimization methods.

# Python

---

# Unconstrained Optimization using `scipy.optimize.minimize`

$$\min_x f(x) = e^{(x_1-2)^2 + (x_2-0.5)^2}$$

```
import numpy as np
import scipy.optimize as optim

# Define the objective function
def fcn(x):
    return np.exp((x[0]-2)**2 + (x[1]-0.5)**2)

# Set the initial point and solve the optimization problem
start = [0, 0]
x = optim.minimize(fcn, start, method='Nelder-Mead')

# The solution is x1=2 and x2=0.5
print(f"The solution is x1={x.x[0]} and x2={x.x[1]}")
```

The solution is  $x_1 = 2$  and  $x_2 = 0.5$ .

# Unconstrained Optimization using `scipy.optimize.minimize`

```
import numpy as np
import scipy.optimize as optim

# Define the objective function
def fcn(x):
    return np.exp((x[0]-2)**2 + (x[1]-0.5)**2)

# Set the initial point and solve the optimization problem
x0 = [0.5, 1.5]
options = {'disp': True, 'maxiter': 1000}
x = optim.minimize(fcn, x0, method='BFGS', options=options)
```

- `'disp': True` shows optimization progress at each iteration



# Unconstrained Optimization using `scipy.optimize.minimize`

```
import numpy as np
import scipy.optimize as optim

# Define the objective function
def fcn(x):
    return np.exp((x[0]-2)**2 + (x[1]-0.5)**2)

# Set the initial point and solve the optimization problem
x0 = [0.5, 1.5]
options = {'disp': True, 'maxiter': 1000}
x = optim.minimize(fcn, x0, method='BFGS', options=options)
```

- `'disp':True` shows optimization progress at each iteration
- `'maxiter'` sets the maximum number of iterations

# Unconstrained Optimization using `scipy.optimize.minimize`

```
import numpy as np
import scipy.optimize as optim

# Define the objective function
def fcn(x):
    return np.exp((x[0]-2)**2 + (x[1]-0.5)**2)

# Set the initial point and solve the optimization problem
x0 = [0.5, 1.5]
options = {'disp': True, 'maxiter': 1000}
x = optim.minimize(fcn, x0, method='BFGS', options=options)
```

- `'disp':True` shows optimization progress at each iteration
- `'maxiter'` sets the maximum number of iterations
- `method='BFGS'` selects a Quasi-Newton optimization algorithm

# Unconstrained vs. constrained

```
import numpy as np
import scipy.optimize as optim

# Define the objective function
def fcn(x):
    return np.sum((x - [5, -7, 3, 0, -11])**2)

# Set the initial point
x0 = 0.5
```

Unconstrained:

```
x = optim.minimize(fcn, x0)
```

Enforce  $x \geq 0$  with LinearConstraint:

```
x = minimize(fcn, x0,
             constraints=optim.LinearConstraint(1, 0))
```

Note that  $A = 1$  and  $b = 0$ , so we enforce that  $x \geq 0$